

Task-Conditioned Reinforcement Learning (TC-RL)

Definition 1 Given a Markov Decision Process (MDP) $\mathcal{M} = \langle S, A, P, \iota \rangle$ and a set of tasks Ψ with a prior distribution $\iota_\Psi \in \Delta(\Psi)$, a **task-conditioned policy** is a mapping $\pi : S \times \Psi \rightarrow \Delta(A)$. The **Task-Conditioned Reinforcement Learning (TC-RL)** problem is to find a policy π maximizing the probability of satisfying a task $\psi \sim \iota_\Psi$ by navigating the underlying environment $\mathcal{M}$, i.e.,

$$J(\pi) \triangleq \mathbb{P}_{\substack{\psi \sim \iota_\Psi \\ \tau \sim \mathcal{M}, \pi, \psi}} [\tau \models \psi], \quad (1)$$

where τ is a trace generated by conditioning π on ψ and running it in $\mathcal{M}$ and $\tau \models \psi$ denotes that the generated behavior τ satisfies the given task ψ . The objective of TC-RL is to solve for the optimal policy $\pi^* \in \arg \max_{\pi} J(\pi)$, maximizing the probability of satisfying $\psi \sim \iota_\Psi$.

- **Automata-conditioned RL $\rightarrow$ condition on**

- **Instruction-conditioned RL $\rightarrow$ condition on**

RAD Embeddings



Qwen3-Embedding



condition

• **Automata-conditioned RL**

• **Instruction-conditioned RL**

condition





RAP Embeddings

